



# OptalCP

Constraint Programming with Parallel Search and Reinforcement Learning-Based Acceleration

Petr Vilím · ScheduleOpt

Vilém Heinz · Czech Technical University in Prague

Scheduling Seminar · [schedulingseminar.com](http://schedulingseminar.com)



# What is OptalCP?

- Constraint Programming solver for scheduling problems.
- From the outside similar to *IBM ILOG CP Optimizer*.
  - Similar modeling language and concepts.
  - Interval variables, sequences, cumulative resources.
- From the inside, completely different.
  - Modern architecture, designed for **parallel search**.
  - Written in **C++20**, APIs in **TypeScript/JavaScript** and **Python**.

Today's focus: How does the solver work inside?



# What makes OptalCP Different?

I've built CP solvers before. Now I'm free to rethink **EVERYTHING**.  
In particular the **internals**.

## Architecture:

- Built for speed from the ground up.
- True parallelism.
- Heterogeneous workers.
- External heuristic hybridization.

## Modeling & API:

- Native **Python** and **TypeScript** APIs.
- **Async** event-driven solving.
- Integers with **optional presence**.
- New modeling constructs.

It's not "just faster" — it's a different architecture that enables new capabilities.



# What makes OptalCP Different?

I've built CP solvers before. Now I'm free to rethink **EVERYTHING**.  
In particular the **internals**.

## Architecture:

- Built for speed from the ground up.
- True parallelism.
- Heterogeneous workers.
- External heuristic hybridization.

## Modeling & API:

- Native **Python** and **TypeScript** APIs.
- **Async** event-driven solving.
- Integers with **optional presence**.
- New modeling constructs.

It's not "just faster" — it's a different architecture that enables new capabilities.

For academic licenses, send me your GitHub username.



# Inside the Solver

## Propagation

Remove infeasibilities

## LNS

Large Neighborhood  
Search

## FDS

Failure-Directed Search

## FDS Dual

Failure-Directed Search  
Dual

Every algorithm has **strengths** and **weaknesses**.



# Propagation

## Propagation

*Always in the Party*

**Role:** Support

**Action:** Remove impossible values

**Produces:** Smaller domains

✓ Prunes domains   Detects infeasibility

✗ Can't solve alone



# Propagation Algorithms for Scheduling

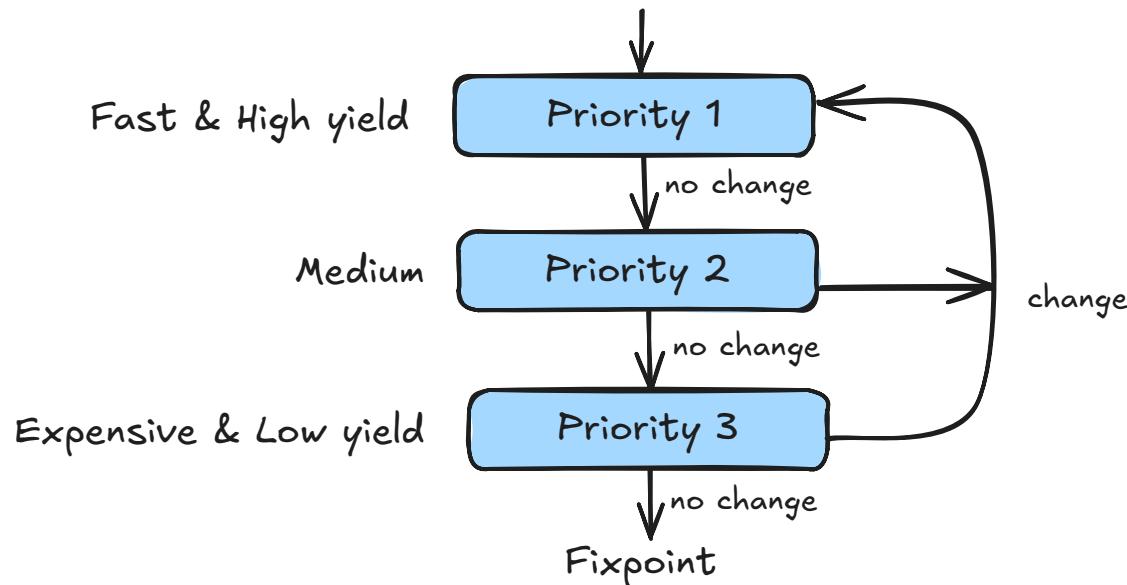
Algorithms for resource constraints in OptalCP:

- Detectable Precedences
- Edge-finding
- Not-first / Not-last
- Timetabling
- Timetable Edge-Finding



# Propagation

In every search node, **propagation** removes impossible values until fixpoint or infeasibility.





# Large Neighborhood Search (LNS)

## LNS

*Large Neighborhood Search*

**Type:** Local search with CP repair

**Assumes:** Better solution exists nearby

**Strategy:** Destroy part, repair with CP

**Produces:** Better solutions

✓ Fast solutions      Exploits structure

✗ Requires initial solution

✗ Local optima      No optimality proof



# How *Standard LNS* Works

Suppose the following solution:





# How *Standard LNS* Works

We relax part of it:





# How Standard LNS Works

We solve the relaxed problem:



**The problem:** Standard LNS fixes the *values* of non-relaxed variables.

Since these are *times*, we can't improve the makespan unless we relax more variables.

Keeping variable values fixed is too restrictive.



# How Standard LNS Works

We solve the relaxed problem:



**The problem:** Standard LNS fixes the *values* of non-relaxed variables.

Since these are *times*, we can't improve the makespan unless we relax more variables.

Keeping variable **values** fixed is too restrictive.

Idea: Modify relaxation to capture the *structure* of the solution instead.

*Philippe Laborie, Daniel Godard:*

Self-adapting large neighborhood search: Application to single-mode scheduling problems



# Partial Order Schedule (POS)

Suppose we have the following solution:





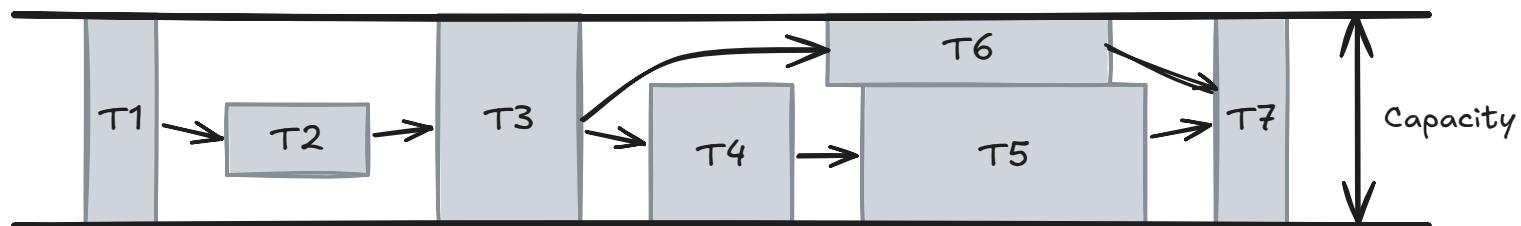
# Partial Order Schedule (POS)

Suppose we have the following solution:



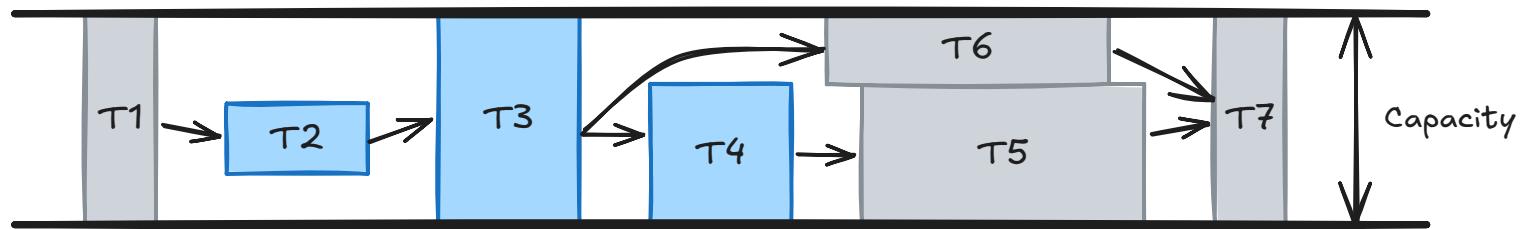
POS = **structure of the solution** = a set of precedences between tasks.

If the variables respect the precedences, resource constraints are automatically satisfied.





# Relaxing with POS




To relax a task, we remove all its precedences.

Transitive precedences between the remaining tasks are added instead.




# Relaxing with POS



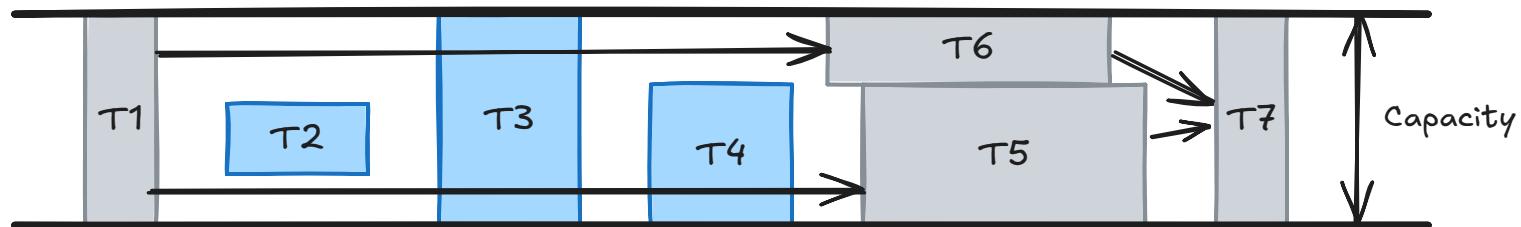
To relax a task, we remove all its precedences.

Transitive precedences between the remaining tasks are added instead.





# Solving with POS



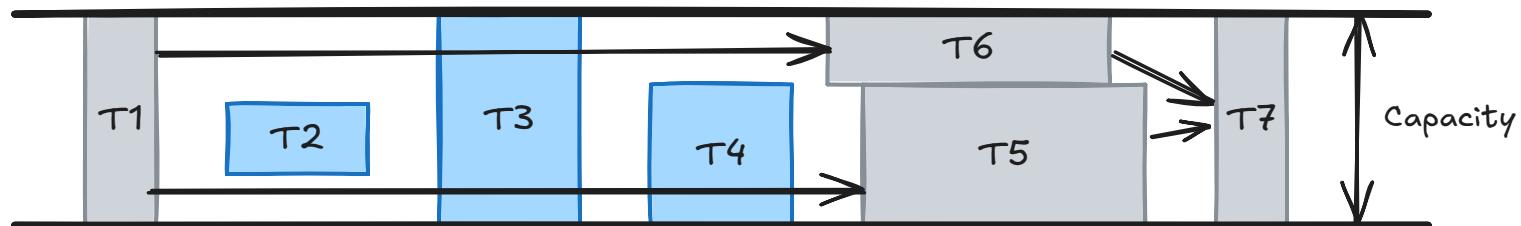
LNS sub-problem:

- Has the same variables (and domains).
- But **more constraints** (precedences from relaxed POS).

Sub-problem solution:



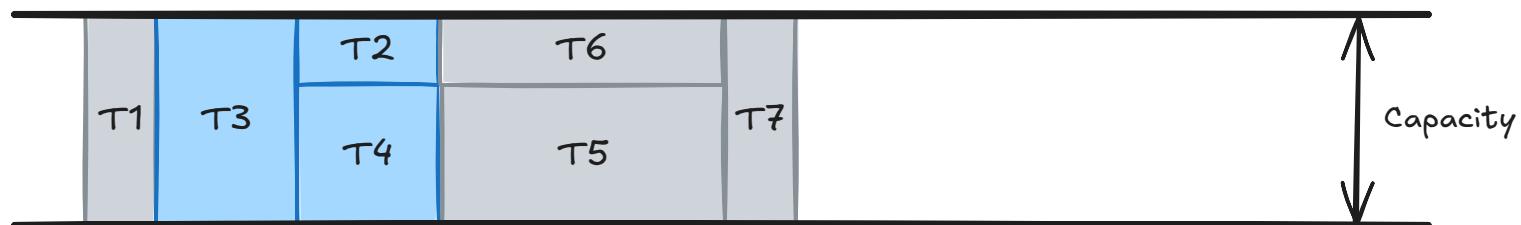
# Solving with POS



LNS sub-problem:

- Has the same variables (and domains).
- But **more constraints** (precedences from relaxed POS).

Sub-problem solution:






# LNS Iterations

Time →

Heuristics, FDS

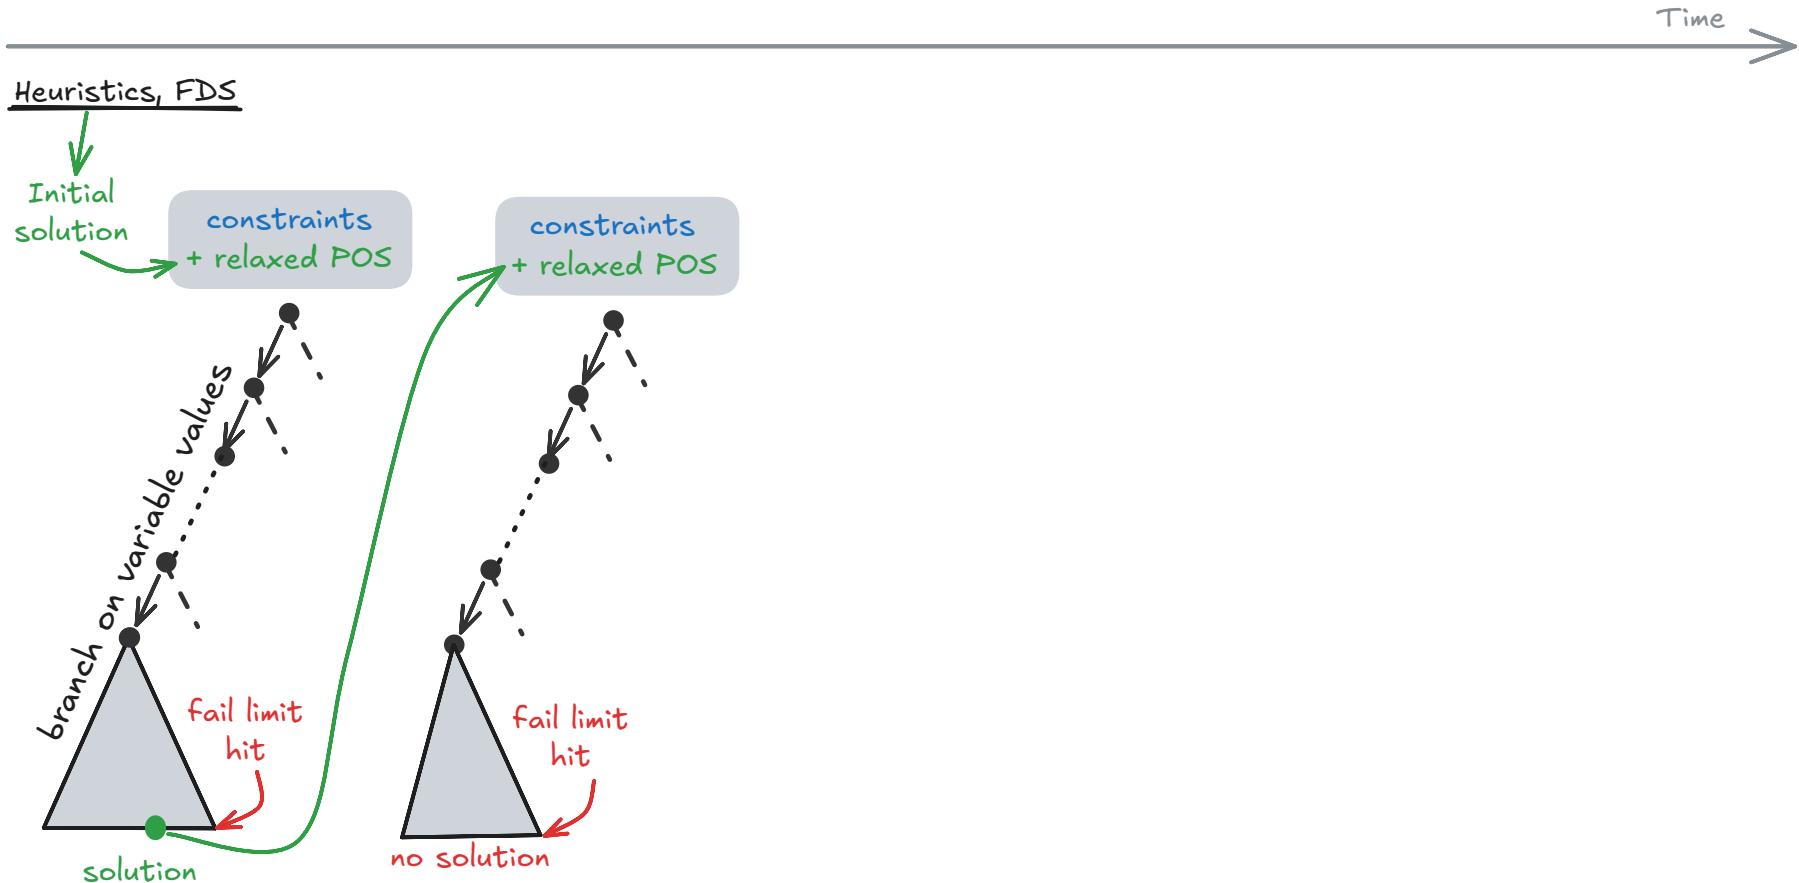
Initial  
solution

constraints  
+ relaxed POS

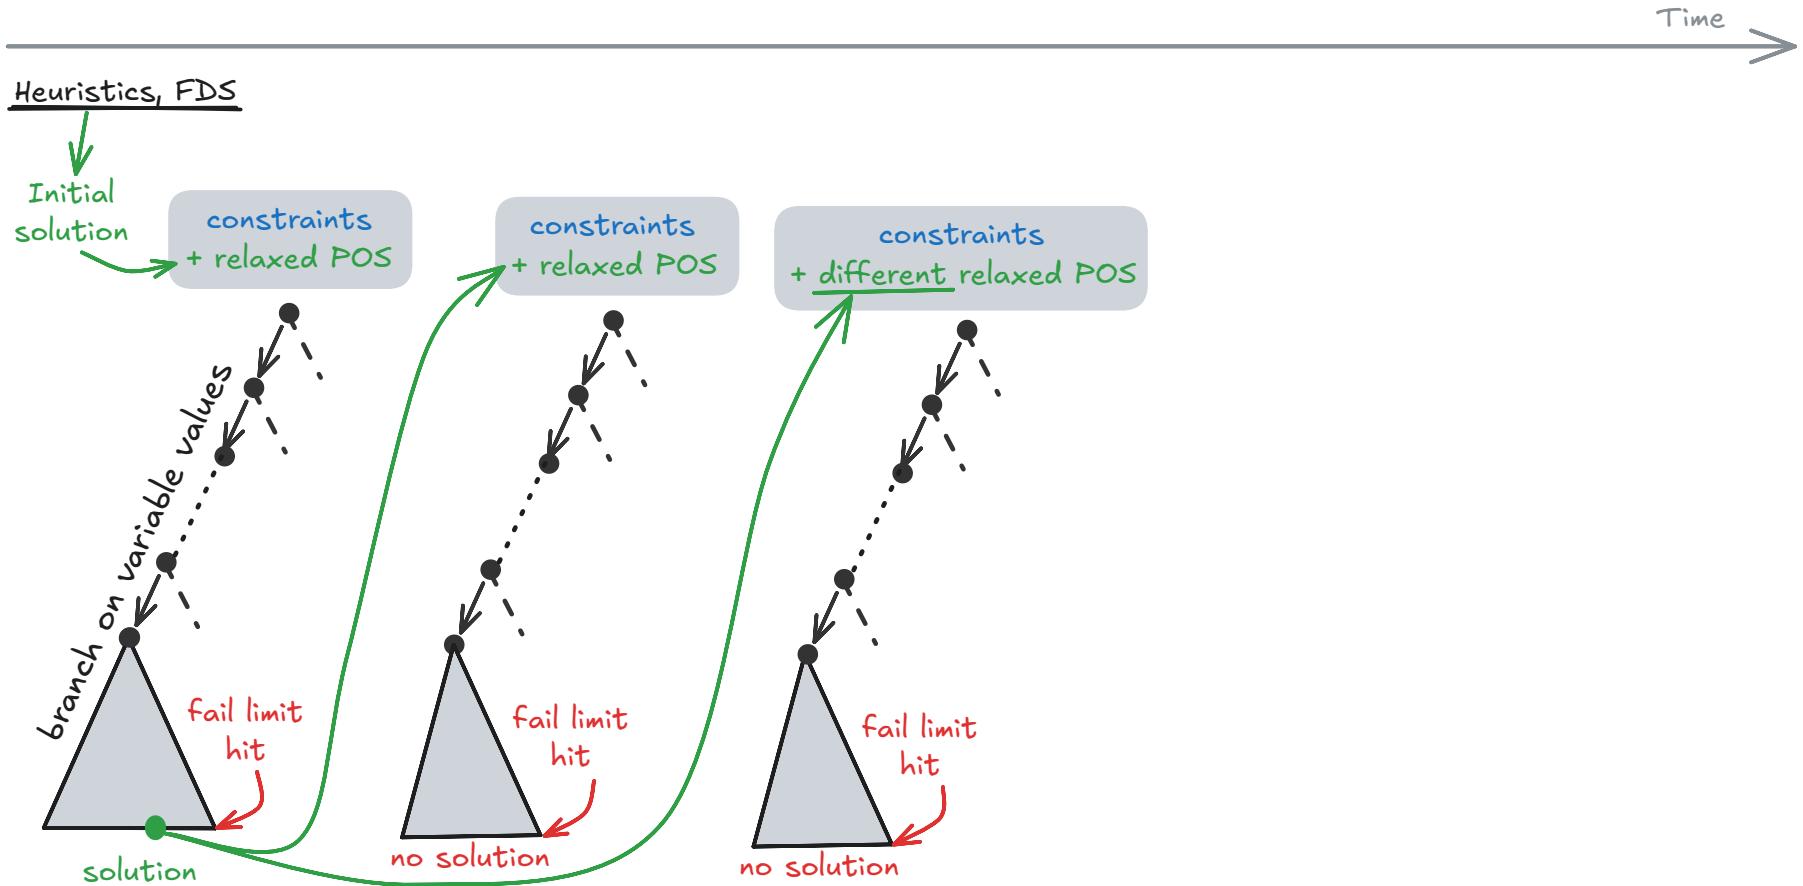




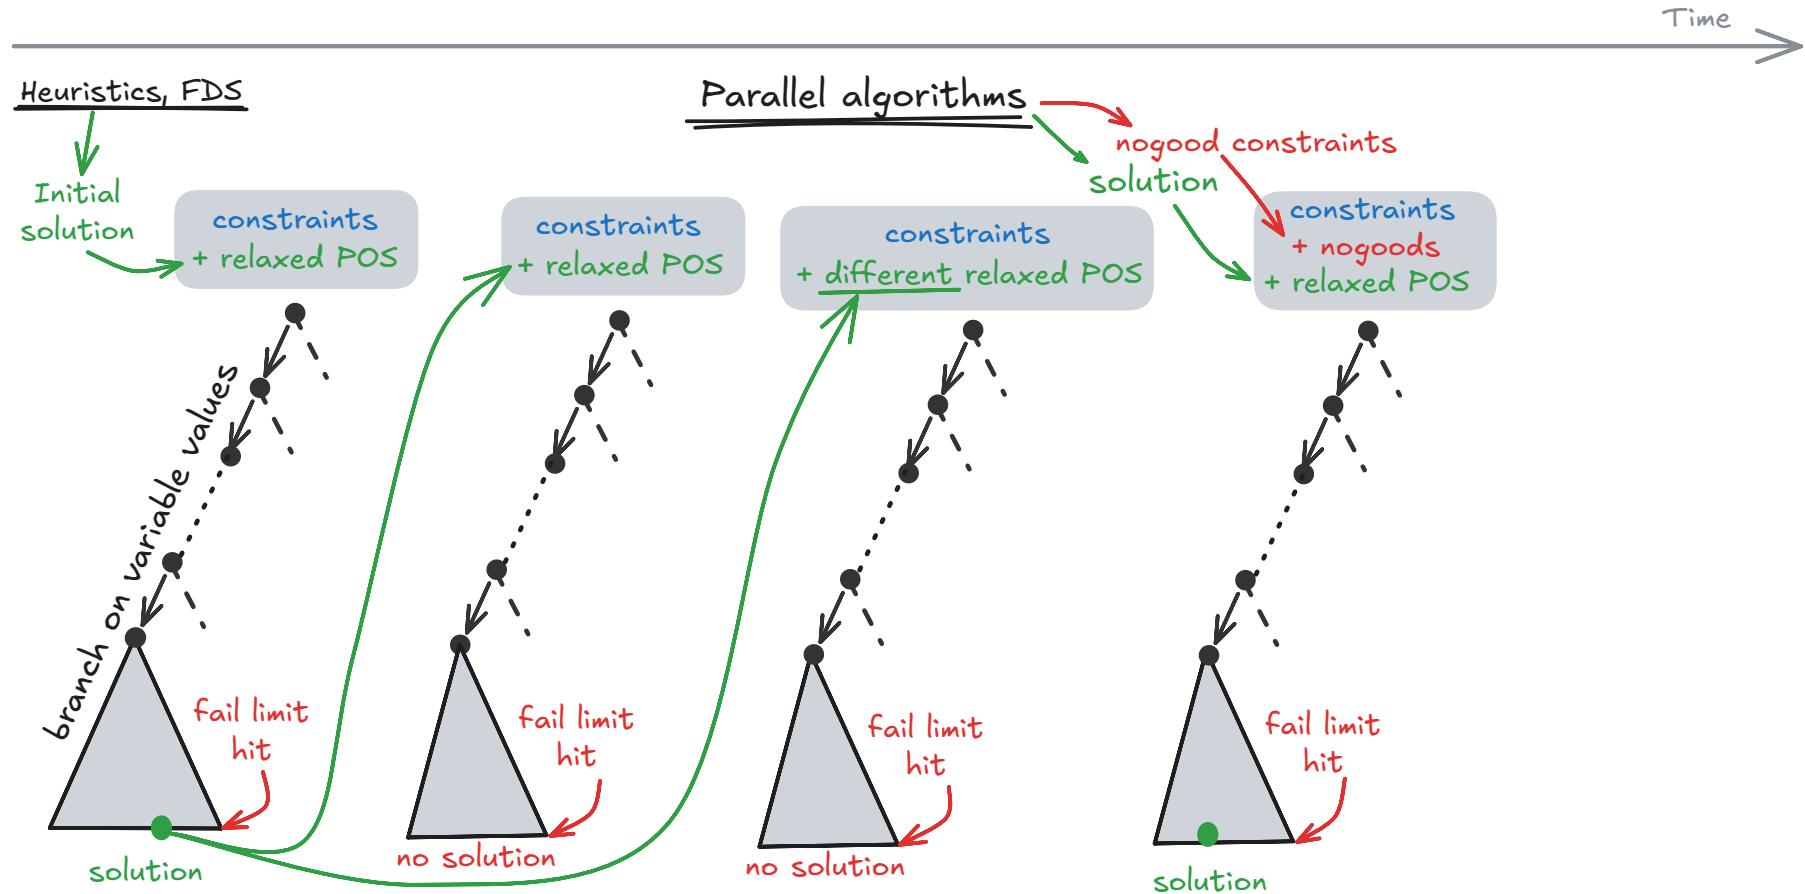
# LNS Iterations



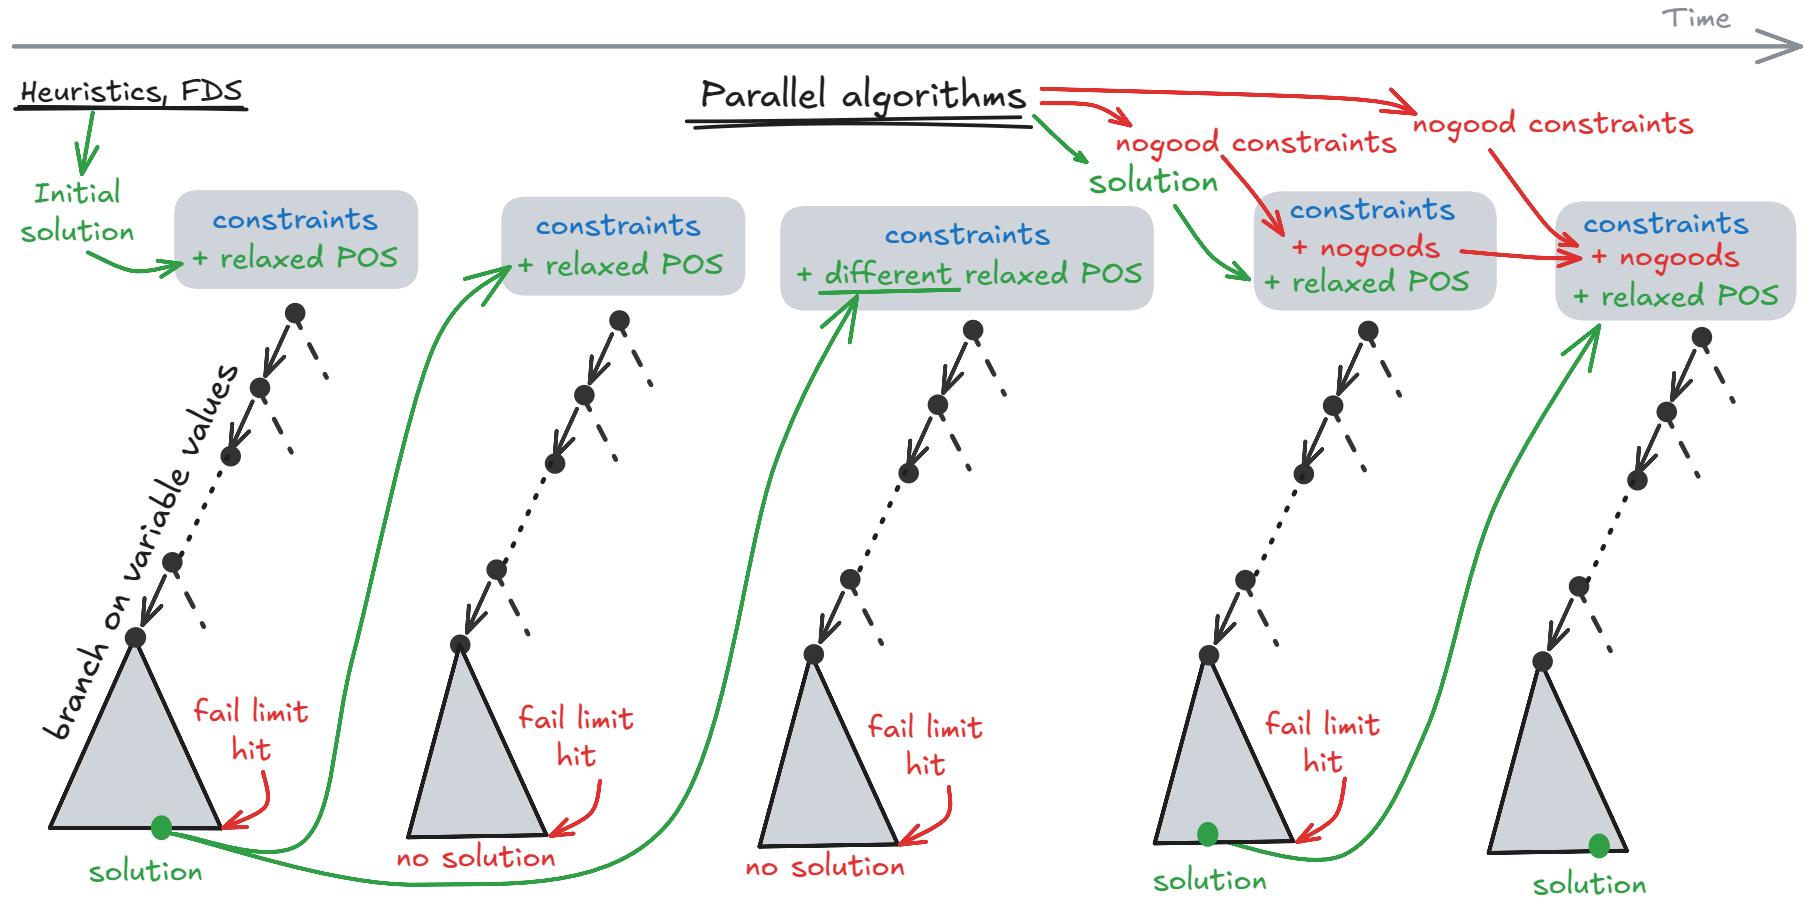
# LNS Iterations



# LNS Iterations



# LNS Iterations





# Failure-Directed Search (FDS)

## FDS

*Failure-Directed Search*

Type: Systematic tree search

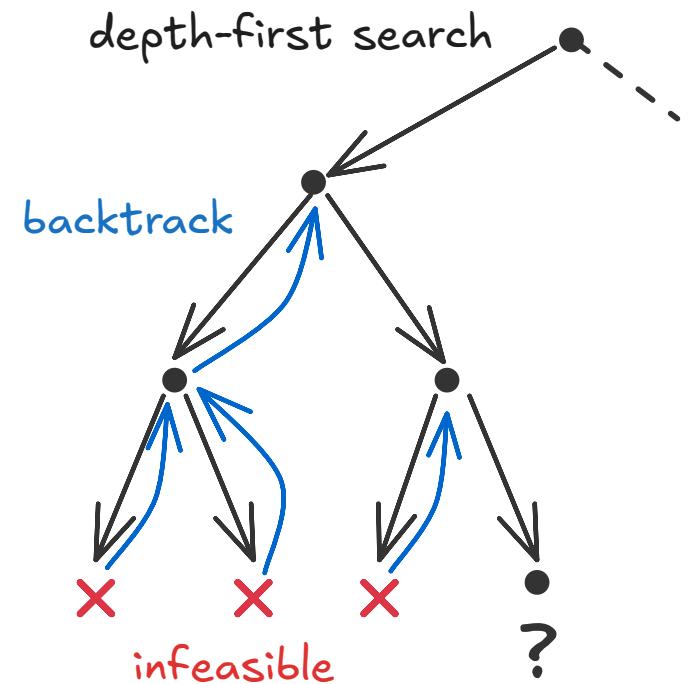
Assumes: Problem is infeasible/hard

Strategy: Learn from failures, restart

Produces: Solutions, proofs, nogoods

✓ Optimality proofs   Lower bounds

✗ Slow   Solutions are a byproduct




Vilém Heinz, Petr Vilím, Zdeněk Hanzálek:

Reinforcement Learning for Search Tree Size Minimization in Constraint Programming:  
New Results on Scheduling Benchmarks



# Simplified FDS


FDS maintains **ratings** on *choices*:





# Simplified FDS


FDS maintains **ratings** on *choices*:





# Simplified FDS


FDS maintains **ratings** on choices:





# Simplified FDS

FDS maintains **ratings** on choices:



$$\text{rating(choice)} := \alpha \cdot \text{rating(choice)} + (1 - \alpha) \cdot \text{penalty}$$



# The FDS - MAB Connection

- MAB algorithms minimize sum of penalties.
- In (simplified) FDS, sum of penalties **is the tree size!**

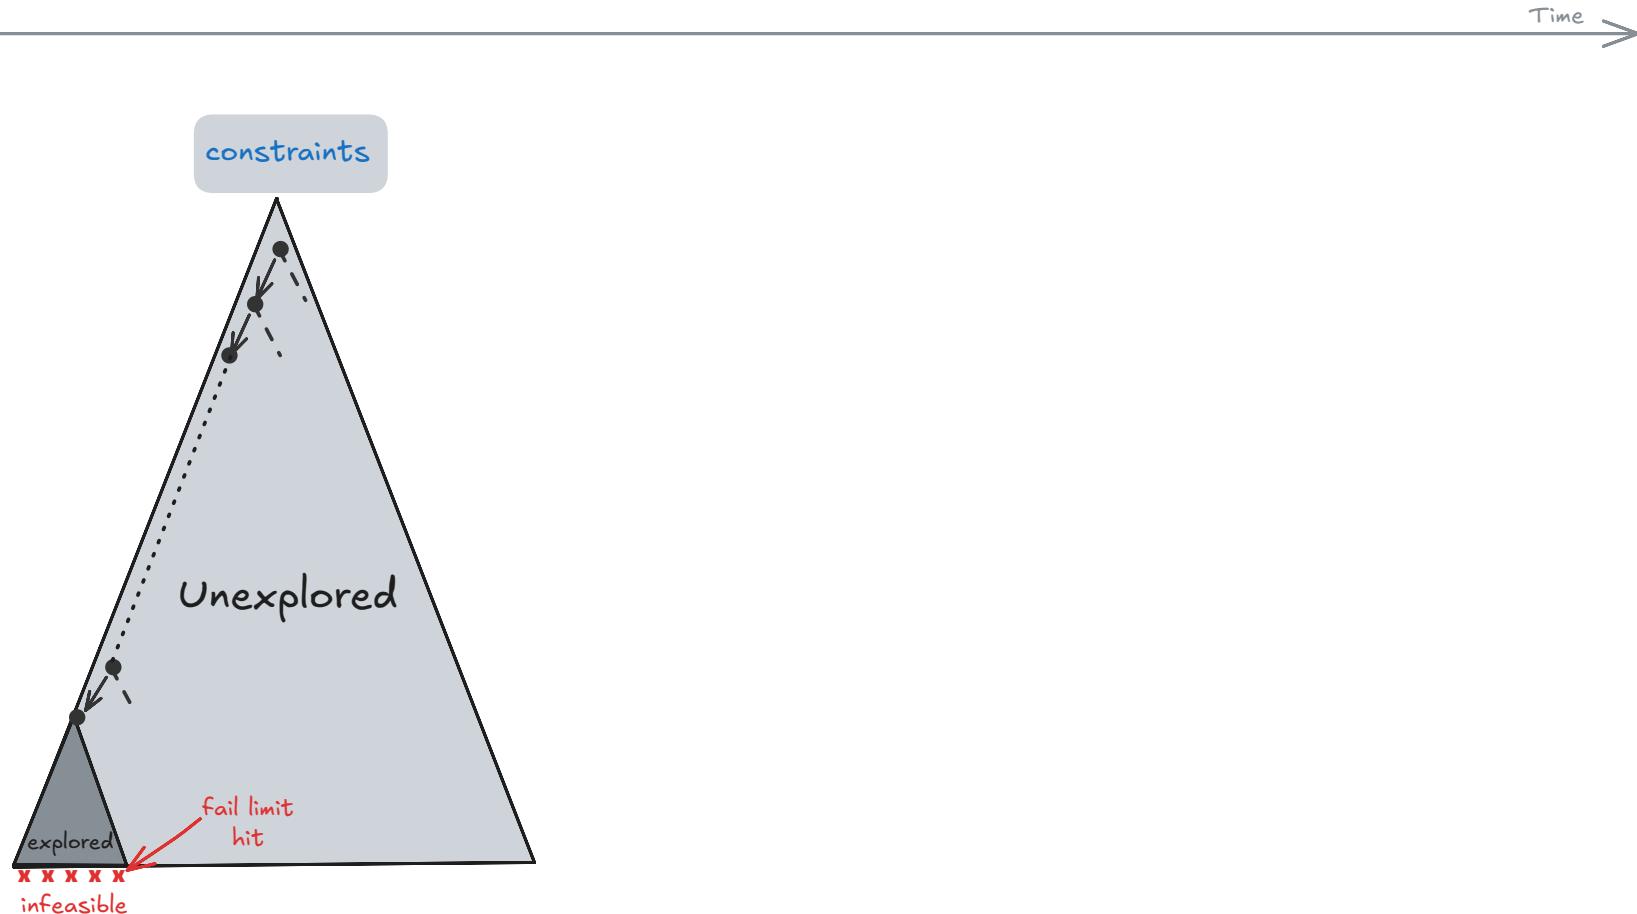
FDS learns to minimize tree size. By design.

I didn't know this when we designed FDS. But it makes sense now.

*Vilém will explore the MAB perspective in more depth.*

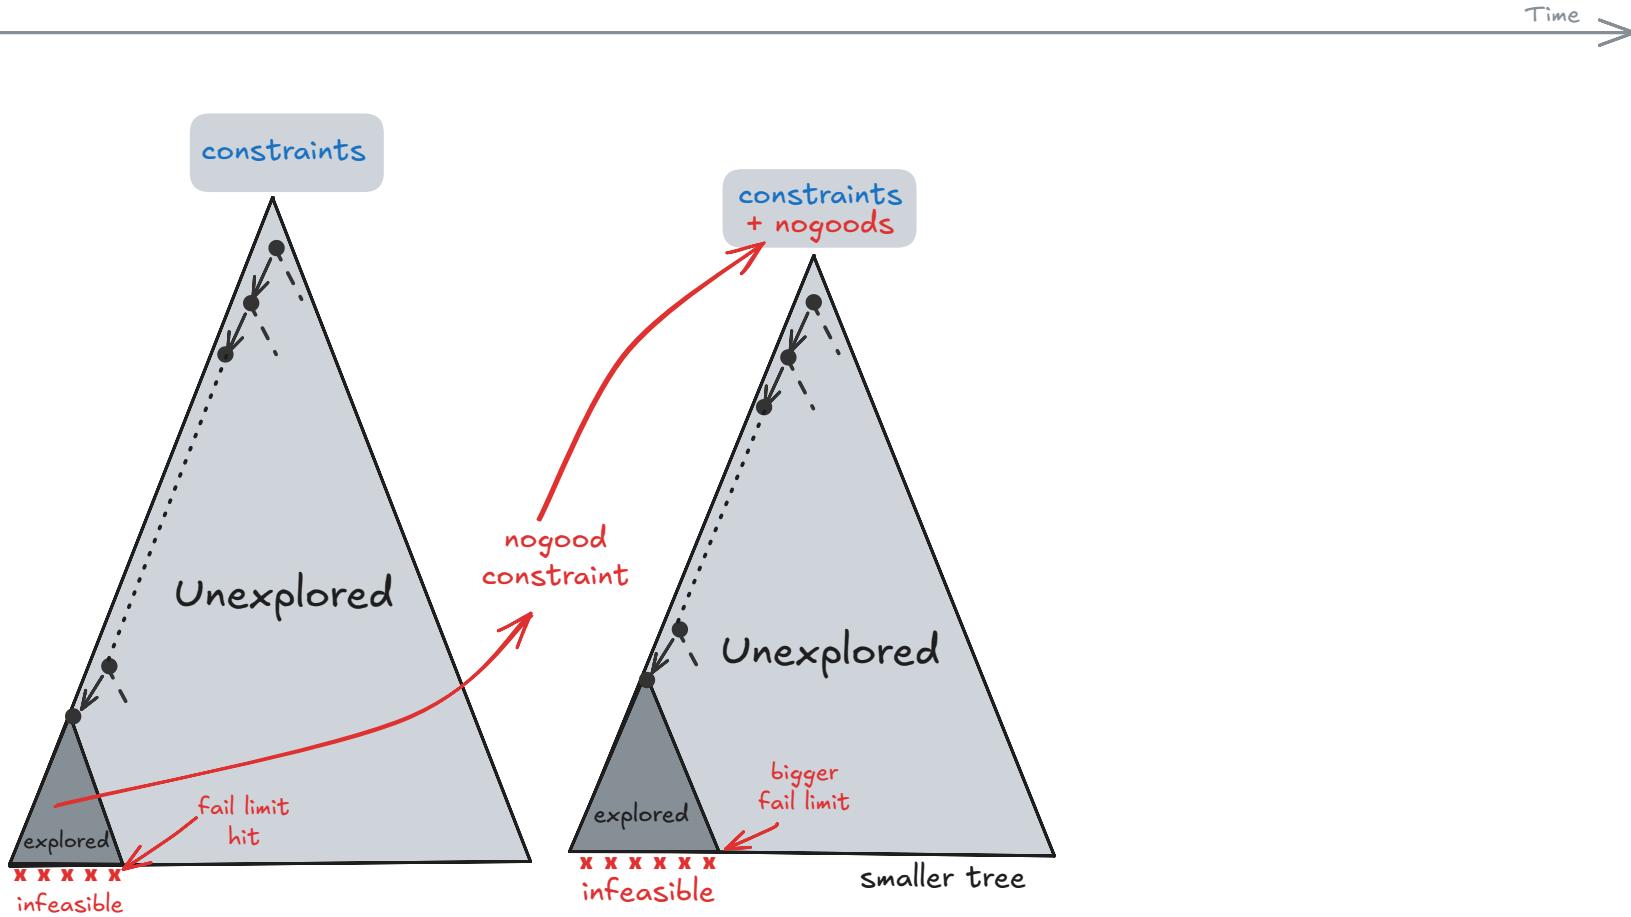


# FDS Restarts






# FDS Restarts



Underlying trees get smaller due to better choices and accumulated nogoods.  
Explored subtrees get bigger due to increased fail limits.



# FDS Restarts



Underlying trees get smaller due to better choices and accumulated nogoods.  
Explored subtrees get bigger due to increased fail limits.



# FDS Dual

## FDS Dual

*Lower Bound Prover*

**Type:** Bound-focused search

**Assumes:** Lower bound can be improved

**Strategy:** Prove infeasible, increment

**Produces:** Tighter lower bound, nogoods

✓ Fast bound proofs    Efficient for LB

✗ No solutions    Not good team player

Smarter version of destructive lower bounds.



# FDS Dual Search

## Efficient lower bound proving

Start with a tight bound, prove it infeasible, then relax.

```
bound = current_LB
while solve(objective <= bound) == INFEASIBLE:
    reportLB(bound + 1)          # Assuming integer
    bound += new_bound_to_try(..) # by parameter FDSDualStrategy
```

FDS Dual workers focus specifically on tightening the lower bound.



# The Perfect Combo

**LNS**

Fast solutions

**FDS**

Proves optimality

**FDS Dual**

Tightens bounds

They complement each other:

**LNS** finds solutions → **FDS** has a better bound

**FDS** finds solutions → **LNS** escapes local optima

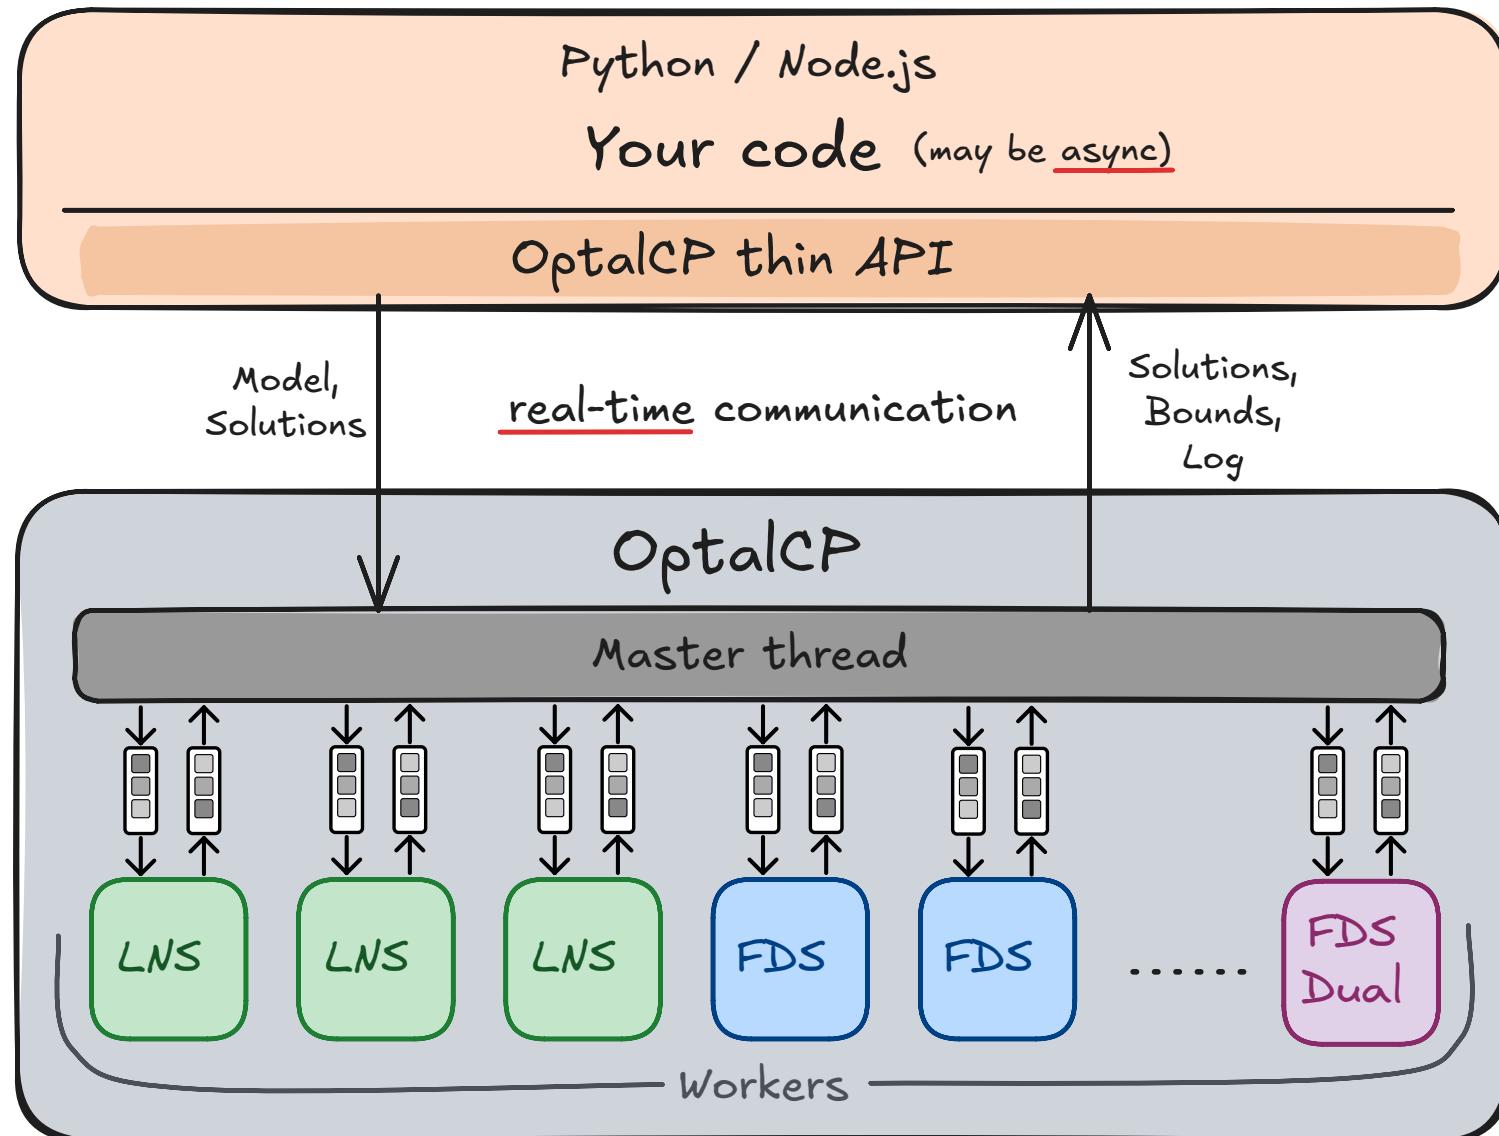
**FDS Dual** proves bounds → Gap shrinks from below

**FDS** **FDS Dual** restarts → Generate nogoods for **LNS** **FDS** **FDS Dual**

**LNS** **FDS** **FDS Dual** prove optimality together → Search ends



# Parallel Architecture





# Heterogeneous Workers

Each worker can be configured independently:

```
model.solve({  
    nbWorkers: 4,  
    workers: [  
        // Fast exploration:  
        { searchType: "LNS", noOverlapPropagationLevel: 2 },  
        // Stronger reasoning:  
        { searchType: "LNS", noOverlapPropagationLevel: 4 },  
        // Optimality focus, escape local optima:  
        { searchType: "FDS", noOverlapPropagationLevel: 4 },  
        // Prove lower bounds:  
        { searchType: "FDSDual", noOverlapPropagationLevel: 4 }  
    ]});
```

Or you can just let OptalCP to decide.



# Hybrid Solution Using Your Algorithm

## Your Algorithm

*Your Secret Weapon*

**Type:** Revolutionary

**Solutions:** Always the best

**Speed:** Blazing fast

**Code:** Beautiful and bug-free

✓ Perfect for the problem

✓ Scales effortlessly

✗ Not in OptalCP   No lower bounds



# Hybrid Solution Using Your Algorithm

## Your Algorithm

*Your Secret Weapon*

**Type:** Revolutionary

**Solutions:** Always the best

**Speed:** Blazing fast

**Code:** Beautiful and bug-free

✓ Perfect for the problem

✓ Scales effortlessly

✗ Not in OptalCP   No lower bounds

### Why hybridize?

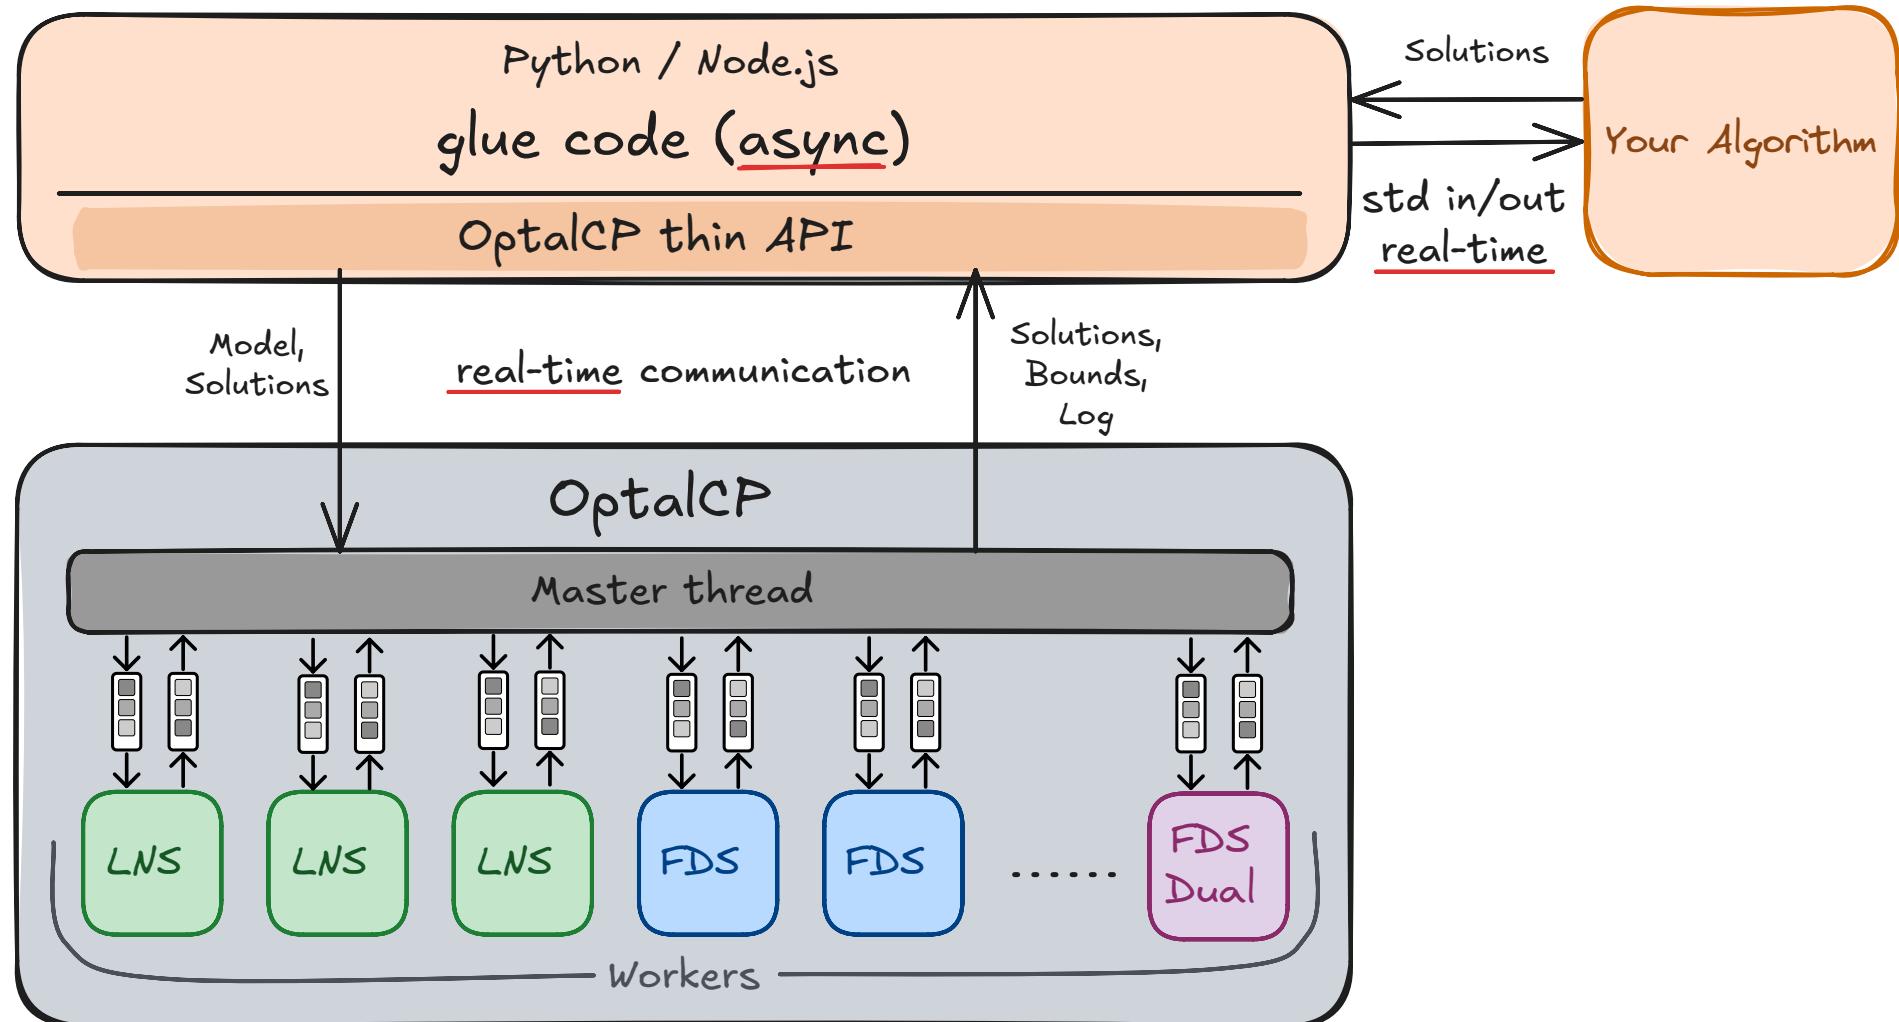
- Solution exchange both ways
- Improved robustness
- Escape local optima
- Better than parallel alone
- Adds optimality gap, stops at optimum

### How to plug in:

- Your algorithm in any language
- Communicates via stdin/stdout
- Short glue code in Python/TypeScript
- Example on GitHub



# Architecture Enabling Hybrid Solution





# Plugging In Your Algorithm

```
your_algorithm = await asyncio.create_subprocess_exec(...)  
solver = cp.Solver()  
  
def on_optalcp_solution(event: cp.SolutionEvent) -> None:  
    serialized = your_solution_format(event.solution)  
    your_algorithm.stdin.write(serialized + b'\n')  
  
async def read_your_solutions() -> None:  
    while True:  
        line = await your_algorithm.stdout.readline()  
        solution = to_optalcp_solution(line)  
        await solver.send_solution(solution)  
  
solver.on_solution = on_optalcp_solution  
asyncio.create_task(read_your_solutions())  
await solver.solve(model, parameters)  
your_algorithm.kill()
```

# Live Demo

demo





# Research Results

- Hybridization with (Meta)heuristics
- Search Acceleration using Reinforcement Learning

Vilém Heinz



# Hybridization with (Meta)heuristics

## Experiments on Scheduling and Routing Problems

Research



# Motivation


Capability	CP Solver	(Meta)heuristics
Bounds and optimality proofs	✓	
Systematic and complete search	✓	
Can prove infeasibility	✓	
Scales to large instances	(✓)	✓
Good anytime behavior		✓
Problem-aware search		(✓)

**Question:** Can we benefit from their complementary nature?

# Goals



- Early search stage:
  - Heuristics provide good feasible solutions
  - Heuristics guide solver's search to promising regions early
- Advanced search stage:
  - Solver incrementally improves and provides bounds
  - Solver helps heuristics to escape local optima
- Overall robustness:
  - Adversarial instances to one method solved by others






# Benchmark Problem Classes

- Two application domains
- Scheduling
  - Flow Shop (FSSP)
  - Job Shop (JSSP)
  - Resource-Constrained Project Scheduling Problem (RCPSP)
- Routing
  - Travelling Salesman Problem (TSP)
  - Capacitated Vehicle Routing Problem (CVRP)
  - Vehicle Routing Problem with Time Windows (VRP-TW)



# Configuration Scheduling


- 5 different metaheuristics for scheduling problems
  - Genetic Algorithm (GA)
  - Large Neighborhood Search (LNS)
  - Particle Swarm Optimization (PSO)
  - Simulated Annealing (SA)
  - Tabu Search (TS)
- 12 threads
  - one thread for each heuristic (5 threads total)
  - rest for solver (5 LNS workers + 2 FDS workers)
- 120s runtime





# Configuration Routing


- 2/3 different (meta)heuristics for routing problems
  - Genetic Algorithm (GA)
  - Local Search (LS)
  - Memetic Algorithm (MA) - only for VRP-TW
- 12 threads
  - one thread for each heuristic (2/3 threads total)
  - rest for solver (7/8 LNS workers + 2 FDS workers)
- 120s runtime





# Practical Example

- Job Shop instance cscmax\_40\_15\_7









- Produces better solution than solver or heuristic portfolio alone
  - Different methods clearly profit from real-time exchange
  - Heuristics are useful again once local optima is escaped



# Results: Instance Solutions

- Green bars denote instances where hybrid outperformed solver, red denote the opposite





# Conclusion

Hybridization improves overall solution quality, anytime behavior and robustness.

- Clear improvements on 3 problem classes
- Overall **60%** of solutions **improved**, only **10% worsened** (mostly slightly) **using same resources**
- Improved overall robustness
  - A few instances unsolved by OptalCP alone were solved by heuristics
- (Meta)heuristics used were not state-of-the-art
  - Still potential for improvements in problems where solver is strong (JSSP, RCPSP)

Problem	Improvement
FSSP	+3% average
JSSP	marginal
RCPSP	marginal
TSP	+6% average
CVRP	marginal
VRP-TW	+11% average




# Accelerating FDS with Reinforcement Learning

## Application of Multi-Armed Bandit Algorithms

Research



# Motivation



- FDS always picks (undecided) choice with best rating
- Good choices with bad initial/recent performance are ignored
  - Choice success can depend on current search state (previous choices)

Question: How to prevent missing such good choices?



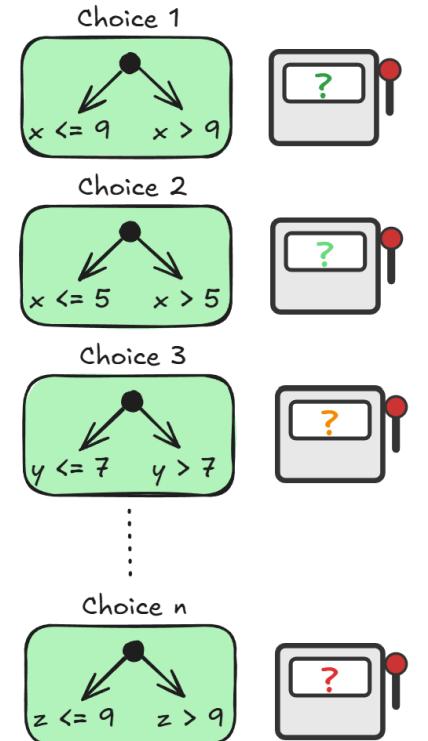
# Goals

- We need to sufficiently test all choices to get an accurate assessment
  - Test all choices initially
  - Revisit bad choices occasionally
- **Enforcing initial choice exploration**
  - Initialize all choices with good rating (optimistic initialization)
- **Reassess choice quality efficiently**
  - FDS choice selection problem  $\approx$  MAB problem



# Reinforcement Learning: MAB Problem

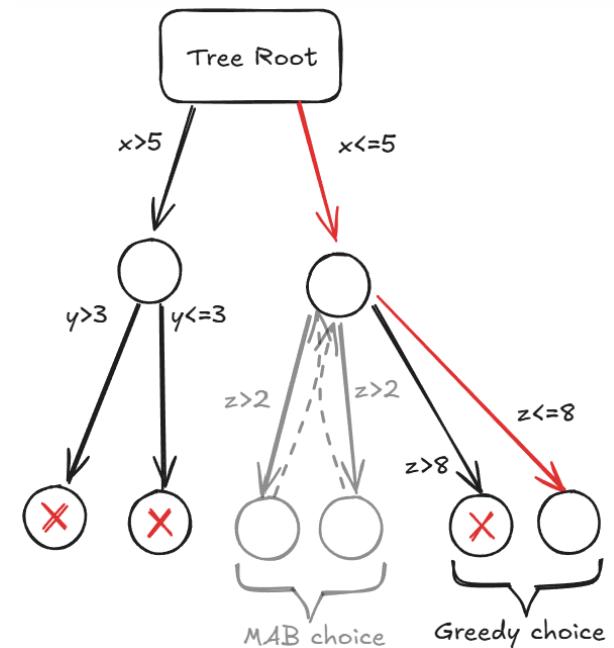
- Exploration-Exploitation dilemma (the problem we have)
  - When to pick the best-known action (exploit)
  - When to test new/under-used actions (explore)
- Multi-Armed Bandit problem
  - Framework solving exploration-exploitation dilemma
  - Different algorithms/ways to handle exploration
  - Epsilon-greedy ( $\epsilon$ ), UCB-1 (U), Boltzmann exploration (B), Thompson sampling (T)
  - MAB reward maximization → Search tree size minimization






# Exploration in FDS Setting

- MAB-based exploration can be costly in FDS
  - Bad exploratory choice = doubling the tree size
- Switch between pure exploitation and MAB strategy
  - In most cases, we exploit
- MAB-based choice rollback
  - “Test run” to evaluate effect and update rating (exploration)
  - Choice is used if it does not increase search tree size, else best-rated choice is used (exploitation)





# Results: Selection Strategies




- Measurements on JSSP and RCPSP (percentage denotes MAB choice probability)
- **10% of Boltzmann exploration with Choice rollback performs the best**
- UCB-1 and Thompson embed exploration in action values (actions regain priority)
  - Degrades performance in FDS exploitation-heavy setting



# Conclusion

Application of extended MAB algorithm with optimistic initialization roughly halved the computation time required by FDS in JSSP and RCPSP instances.

- Improved a large number of lower bounds for both problems
  - 78/84 of open standard JobShop (JSSP) instances
  - 226/393 of open standard RCPSP instances
  - A few instances for both problems were closed
  - 900s time limit per instance





# Thank You!

## Questions?

Website: <https://optalcp.com>

Benchmarks GitHub: <https://github.com/scheduleopt/optalcp-benchmarks>

Academic Licenses: Send me your GitHub username